

CS Educators Stack Exchange: Exploring User Interactions

Sukanya Moudgalya & K. Bret Staudt Willet

12/7/2018

Introduction

The Stack Exchange Network is one of the largest question-and-answer platforms in the world, founded in 2009 and boasting over 100 million unique monthly visitors by 2015. It has many subject-specific forums for coding related topics, Mathematics, Engineering, Photography, etc.—133 different sites as of October 2018. More recently, Stack Exchange has recently opened up discussion forums for Math educators (2013) and Computer Science educators (2016) as well. For this project, we will focus on the users of the discussion forum Computer Science Educators’ Stack Exchange (CSEd SE). The description of the forum calls it ‘A question and answer site for those involved in the field of teaching Computer Science.’

In this forum, registered users can ask questions on a wide range of topics related to Computer Science (CS) education. Once they ask a question, they can assign it ‘tags’, so that similar groups of questions can be identified by a common theme or idea. Some tags that are currently present on the forum include ‘curriculum design,’ ‘student motivation,’ ‘algorithm,’ etc. Users can also reply to (i.e., answer) questions posed by other users, comment on both questions and answers, upvote and downvote questions and answers, and so on—through the voting mechanism, users are held accountable to what they post on the forum. The content on the forum is thus regulated by the forum users and members to a large extent.

In addition to ‘normal’ user regulation, a set of registered users also act as ‘moderators’ or ‘editors’. Editors can edit the text of what other people post. Moderators, on the other hand, have more power. They can follow up with posts that are flagged as spam or offensive, delete and manage other users, combine or collapse tags, lock posts so that no one can change them or vote on them, and so on (read more here).

The users on this forum range from actual K-12 and college-level CS educators to working professionals in the software industry, students, and other people interested in CS education. The users also belong to many different countries, such as the USA, UK, Canada, Ireland, India, Finland, etc. Each user has reputation points primarily based on the quality of their posts, such as user-generated votes their posts receive. The reputation changes across time are accessible for people to see.

Purpose

This study explores the kinds of interaction between the users, with a focus on the top voted questions. It seeks to describe the kind of users, based on their profession, location, reputation, etc., that tend to ask questions or respond to the top questions. It also seeks to understand if the users tend to cluster around certain question tags (topic area) or if the clusters are agnostic to the topics and user characteristics.

In the second part of the study, we will focus on trying to interpret the reasons for users choosing to answer the questions. It will explore if the reasons are dependent to the topic tag or some characteristic of the users who ask the question.

Theoretical Framework and Research Questions

Communities of practice are "...groups of people who share a concern or a passion for something they do and learn how to do it better as they interact regularly" (Wenger, 2011, p. 1). These communities have the following traits: (a) a domain of shared interest; (b) people who engage in joint activities and discussions; and (c) members who are practitioners (such as educators). Further, according to Wenger, they should demonstrate (a) mutual engagement (such as a dialogue instead of one directional information flow); (b) a joint enterprise; and (c) a shared repertoire. These three traits represent "...three dimensions of the relation by which practice is the source of coherence of a community" (Wenger, 1998, p. 72).

The CSEd SE has many elements that could potentially make it a 'community of practice'. For instance, there are clear boundaries in terms of membership to the forum. The forum has a unique purpose, that is to build knowledge and dialogue in the domain of CS education. Further, the discussions on the forum are not generally a unidirectional venture. The commenting feature, for instance, allows scope for a back and forth dialogue. Finally, the presence of moderators and editors for maintenance of the forum, scope of democratic elections for the selection of moderators, reputation points based on user-generated voting system, gives CSEd SE the traits of a community of practice. Indeed, a literature review of informal online teacher communities revealed that the communities of practice framework has often been used to study teacher interactions in online forums (Macià & García, 2016).

Using this communities of practice framework, we focused our study to answer three research questions:

- * *RQ1*. What are the characteristics of users who contributed to the CSEd SE's core network of interactions?
- * *RQ2*. Does network clustering occur in the context of the CSEd SE's core network of interactions? If so, what are some characteristics of these clusters?
- * *RQ3*. From all the questions that users are exposed to in the CSEd SE's core network of interactions, what factors predict the questions that users to respond to?

Method

Data Collection

The data from the CSEd SE forum were collected through data mining, using the statistical software *R* (R Core Team, 2018). In particular, we used the *R* package *stackr* (Robinson, 2018) to extract data from the CSEd SE website. The *stackr* package helps obtain the "read-only features of the Stack Exchange API with the ability to download information on questions, answers, users, tags, and other aspects of the site so that they can be analyzed in *R*".

Using *stackr*, we were able to collect the metadata for CSEd SE questions, answers, comments, tags, and users in the week of 7th-13th June 2018. These metadata included information such as question identity number, answer identity number, user identity number, tag name, reputation scores of users, the date they joined the forum, and so on. This data collection resulted in a corpus of 559 questions asked by 210 different users, 2,675 answers contributed by 651 different users, and 7,209 comments from CSEd SE.

Data Analysis

First, we loaded the necessary *R* packages and input our data.

We then reshaped the data for analysis. We used the *R* package *igraph* (Csárdi, 2018) to create a graph of the CSEd SE's *core network of interactions*, focusing on nodes in the network (i.e., contributors) with degree of three or higher. In other words, we excluded users who participated as "tag-alongs," perhaps answering a question or two but not interacting with other users in a significant way.

Results

RQ1. What are the characteristics of users who contributed to the CSEd SE's core network of interactions?

To answer RQ1, we first tallied the number of questions, answers, and contributors to the CSEd SE during our data collection time period.

```
csed_ques$question_id %>% unique %>% length # number of distinct questions in CSEd SE
## [1] 559
csed_ques$display_name %>% unique %>% length # number of total, distinct contributors
## [1] 210
csed_ans$question_id %>% unique %>% length # number of questions answered in our time window
## [1] 553
csed_ans$answer_id %>% unique %>% length # number of answers in our time window
## [1] 2675
csed_ans$display_name %>% unique %>% length # number of people who answered questions in our time window
## [1] 651
## these are the different responders in the core network
csed_full_weighted %>% filter(weight > 0) %>%
  select(responder) %>% unlist %>% unique %>% length
## [1] 175
## these are the different question askers in the core network
csed_full_weighted %>% filter(weight > 0) %>%
  select(asker) %>% unlist %>% unique %>% length
## [1] 164
## these are total, distinct contributors (askers & responders) in the core network
c({csed_full_weighted %>% filter(weight > 0) %>% select(responder) %>% unlist},
  {csed_full_weighted %>% filter(weight > 0) %>% select(asker) %>% unlist}
) %>% unique %>% length
## [1] 257
```

We then calculated descriptive statistics.

```
##          Degree In-Degree Out-Degree
## mean:         9.96      4.98      4.98
## sd:          17.88      9.38     10.35
## median:        5.00      3.00      3.00
## min:           1.00      0.00      0.00
## max:         146.00     80.00     101.00
## skewness:       5.14      4.44      5.36
## kurtosis:      34.06     27.66     40.59
```

RQ2. Does network clustering occur in the context of the CSEd SE’s core network of interactions? If so, what are some characteristics of these clusters?

As a reminder, a quick summary of the the CSEd SE core network of interactions during the time period of our study:

```
##  
## Number of nodes:    257  
## Number of edges:    1300  
## Diameter:           6
```

Clustering: Spinglass Algorithm

A community is a set of nodes with many edges inside the community and few edges between outside it (i.e. between the community itself and the rest of the graph).

The *spinglass clustering algorithm* maps community detection onto finding the ground state of an infinite range spin glass. Csárdi, Nepusz, and Airoldi (2016, pp. 132-133) explained:

The clustering method of Reichardt and Bornholdt (2006) is motivated by spin glass models from statistical physics. Such models are used to describe and explain magnetism at the microscopic scale at finite temperatures. Reichardt and Bornholdt (2006) drew an analogy between spin glass models and the problem of community detection on graphs and proposed an algorithm based on the simulated annealing of the spin glass model to obtain well-defined communities in a graph. A spin glass model consists of a set of particles called spins that are coupled by ferromagnetic or antiferromagnetic bonds. Each spin can be in one of k possible states. The well-known Potts model then defines the total energy of the spin glass with a given spin configuration. . . Spins and interactions in the Potts model are very similar to graphs: each spin in the model corresponds to a vertex, and each interaction corresponds to an edge. . . Reichardt and Bornholdt (2006) gave efficient update rules for the above energy function, making it possible to apply a simulated annealing procedure to find the ground state of the model that corresponds to a low energy configuration. Their algorithm starts from a random configuration of spins and tries to flip all the spins once in each time step. After each individual spin flip, the energy of the new configuration is evaluated.

In other words, the spinglass clustering algorithm partitions the vertices into communities by optimizing an energy function. The initial R code to produce spinglass clusters is straightforward:

```
csg_0 <- csed_graph %>% cluster_spinglass # creates the clusters  
csg_0$membership %>% unique %>% length # number of clusters/communities/groups
```

```
## [1] 8
```

One of the important outcomes of this method is the *modularity* value M . Modularity measures how good the division is, or how separated are the different vertex types from each other. The spinglass algorithm looks for the modularity of the optimal partition. For a given network, the partition with maximum modularity corresponds to the optimal community structure (i.e., a higher M is better).

Note also that if $M = 0$, all nodes belong to one group; if $M < 0$, each node belongs to separate community.

```
csg_0$modularity # The modularity of a graph with respect to some division (or vertex types)
```

```
## [1] 0.2703746
```

Identifying the “Typical” Number of Clusters Returned with the Spinglass Algorithm

It is important to note that a different result is returned each time the spinglass clustering algorithm is run. For this reason, we needed to run a number of simulations to see what the “typical” number of clusters are.

We ran the algorithm 100 times and looked at the mean and median number of clusters obtained. We made a note of a `seed` that produced the median number of clusters, confirmed that this was reproducible, and then set this seed so that all future work will be run with this same clustering configuration.

```
##
## number of tests: 100.00
## mean:          7.51
## sd:            0.64
## min:           6.00
## max:           9.00
## skewness:      -0.72
## kurtosis:       2.80
## median:        8.00

##
## Number of nodes:      257.0000
## Number of edges:      1300.0000
## Number of clusters:    8.0000
## Modularity:           0.2714

## [1] Size of each cluster:
## [1] 30 58 49 11 36 21 27 25
```

Test of Statistical Significance for Spinglass Clusters

The test for statistical significance for spinglass clustering is a bit different than the familiar tests that return p -values (Csárdi, Nepusz, & Airoldi (2016, pp. 132-138).

The idea behind this test of significance is that a random network of equal size and degree distribution as our observed network should have a lower modularity score—that is, if the observed network does in fact have statistically significant clustering.

The following R procedure generates 100 randomized instances of our network (with the same size and degree distribution) using the `sample_degseq()` function. The `method = 'v1'` ensures that there are no loop edges in the randomly generated networks. We then applied the spinglass clustering algorithm to each of the 100 randomized instances of the network.

A ‘0’ result from this procedure indicates that no randomized networks have community structure with a modularity score that is higher than the one obtained from the original, observed network. Hence a ‘0’ result means that our network has significant community structure; any non-zero results means that the detected spinglass clusters are not statistically significant.


```
#degrees <- csed_graph %>% as.undirected %>% degree(mode='all', loops=FALSE)
#qr_vl <- replicate(100, sample_degseq(degrees, method="v1", simplify=FALSE) %>%
#      lapply(cluster_spinglass) %>%
#      sapply(modularity)
#saveRDS(qr_vl, "model-output/qr_vl.rds")
qr_vl <- read_rds("model-output/qr_vl.rds")
sum(qr_vl > csg$modularity) / 100
```

```
## [1] 0
```

Network Visualization with Clusters

Finally, we created a visualization of our network structure, using the color palette generated by our spinglass clustering. Here, we used the Fruchterman-Reingold layout algorithm (`layout = 'fr'`), which is appropriate for large (but still with less than 1,000 nodes), potentially disconnected networks.

```
## IGRAPH c5c35c9 DN-- 257 1300 --
## + attr: name (v/c), degree (v/n), in_degree (v/n), out_degree
## | (v/n), cluster_weight (e/n)
```



RQ3. What factors predict the questions that users to respond to?

Describing Selection Models

Level 1 Selection Model

The selection model at level 1, for responder i , answering question q , asked by user i' :

$$\log\left(\frac{p(\text{Answering a certain question} = 1)}{1 - p(\text{Answering a certain question} = 1)}\right) = \theta_i + \theta_{i'} + \theta_q + \theta_1(\text{Prior relationship of } i \text{ and } i') + \theta_2(\text{Prior connection between } i \text{ and } q)$$

Where

1. θ_i = Responder effect
2. $\theta_{i'}$ = Asker effect
3. θ_q = Question effect
4. Prior relationship of i and i' = 1 if i had answered one of i' 's question or if i' 's had answered one i 's question previously; 0 if i and i' had no such interaction.
5. Prior connection between i and question tag q = 1 if i had previously answered any question with the same topic tag(s) as question q ; 0 if i had no such connection with q .

Level 2 Selection Models

The selection model at level 2, for responder i ;

$$\theta_i = \alpha_0 + \alpha_1(\text{Reputation points}) + \alpha_2(\text{Role}) + \alpha_3(\text{Duration of membership}) + \alpha_4(\text{Number of questions previously answered}) + \epsilon_1$$

Where

1. Reputation points = The numerical reputation points associated with each user.
2. Role = 1 if the responder answering question is an editor or a moderator. 0 otherwise.
3. Duration of Membership = Time in months of being a member of CSEd SE
4. Number of questions previously asked = Total number of questions asked up until present time.

The selection model at level 2, for user i' asking the question;

$$\theta_{i'} = \beta_0 + \beta_1(\text{Reputation points}) + \beta_2(\text{Duration of membership}) + \beta_3(\text{Number of questions previously asked}) + \epsilon_2$$

Where

1. Reputation points = The numerical reputation points associated with each user.
2. Duration of Membership = Time in months of being a member of CSEd SE.
3. Number of questions previously asked = Total number of questions asked up until present time.

The selection model at level 2, for question q ;

$$\theta_q = \gamma_0 + \gamma_1(\text{Number of votes}) + \gamma_2(\text{Number of answers}) + \gamma_3(\text{Number of views}) + \epsilon_3$$

Where


1. Number of votes = Total up-votes the question got on the forum prior to i answering the question.
2. Number of answers = Total answers the question had prior to i answering the question.
3. Number of views = Total views the question got on the forum prior to i answering the question.

Possible sources of multicollinearity

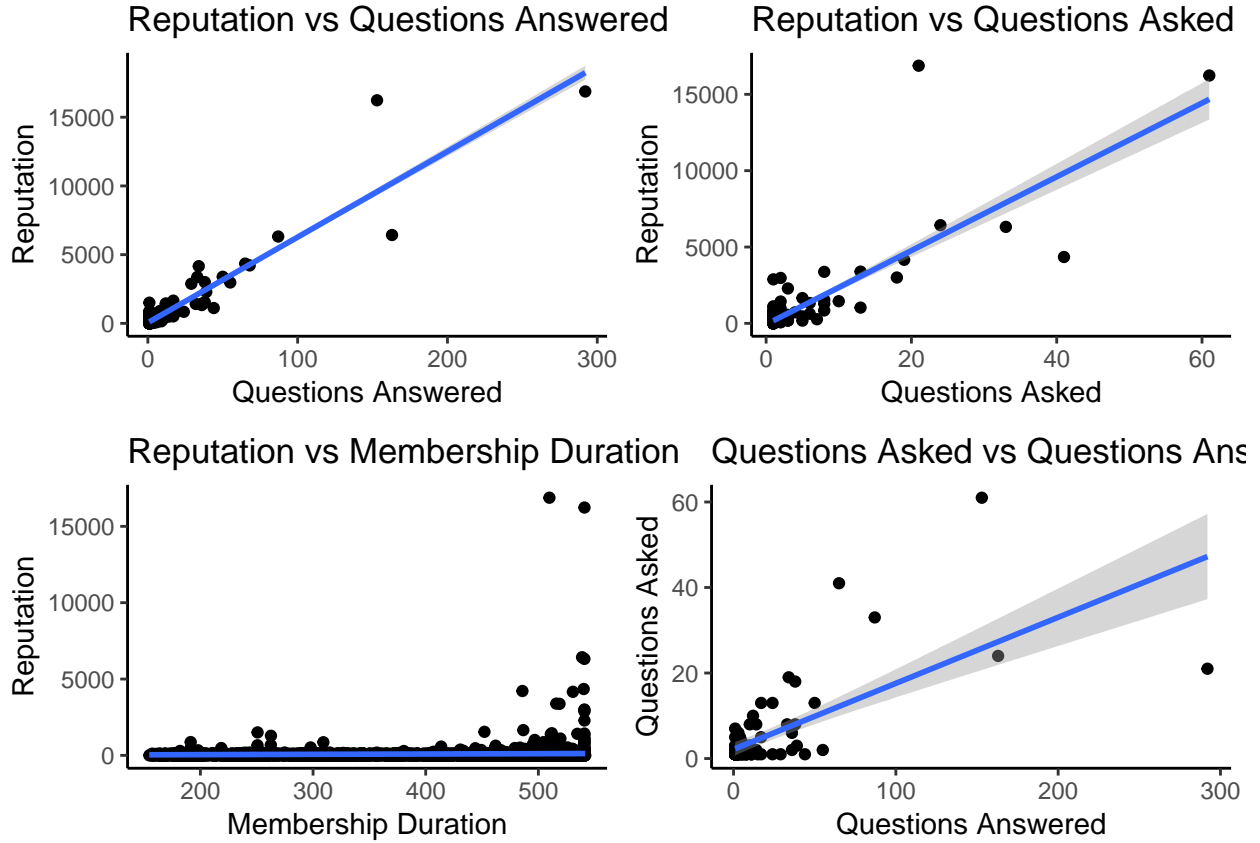
We tested to see if these variables are highly correlated and if we needed to change the models that we proposed. These are the correlations that we tested:

1. Correlations between number of votes, answers, and views
2. Correlations between reputation points and number of questions answered/asked

First, we explored the correlations between the number of votes, answers, and views. Please see the plots and the correlation matrix below. As these variables were highly correlated, we decided to use only one, the number of votes.



```
## [1] Correlations:
##           view_count answer_count score
## view_count      1.00      0.67  0.92
## answer_count    0.67      1.00  0.74
## score           0.92      0.74  1.00
## Warning: Removed 4445 rows containing non-finite values (stat_smooth).
## Warning: Removed 4445 rows containing missing values (geom_point).
## Warning: Removed 4879 rows containing non-finite values (stat_smooth).
## Warning: Removed 4879 rows containing missing values (geom_point).
## Warning: Removed 8 rows containing non-finite values (stat_smooth).
## Warning: Removed 8 rows containing missing values (geom_point).
## Warning: Removed 5000 rows containing non-finite values (stat_smooth).
## Warning: Removed 5000 rows containing missing values (geom_point).
```

[1] Correlations:

```
##               reputation questions_answered questions_asked
## reputation               1.00              0.94             0.81
## questions_answered       0.94              1.00             0.68
## questions_asked          0.81              0.68             1.00
## membership_duration_days 0.18              0.18             0.17
##
##               membership_duration_days
## reputation                   0.18
## questions_answered           0.18
## questions_asked              0.17
## membership_duration_days     1.00
```

UPDATED Level 2 Selection Models

New model at level 2, for responder i ;

$$\theta_i = \alpha_0 + \alpha_1(\text{Reputation points}) + \alpha_2(\text{Role}) + \alpha_3(\text{Duration of membership}) + \epsilon_1$$

(We removed 'number of questions previously answered' as it was highly correlated with reputation points)

New model at level 2, for user i' asking the question;

$$\theta_{i'} = \beta_0 + \beta_1(\text{Reputation points}) + \beta_2(\text{Duration of membership}) + \epsilon_2$$

(We removed 'number of questions previously asked' as it was highly correlated with reputation points)

New model at level 2, for question q ;

$$\theta_q = \gamma_0 + \gamma_1(\text{Number of votes}) + \epsilon_3$$

(We removed ‘number of answers’ and ‘number of views’ as they were highly correlated with number of votes)

Constructing the Selection Models: Linear Mixed-Effects Models

Having conceptually described the selection models of users on CSEd SE choosing to answer certain questions, we then built these in R using *generalized linear mixed-effects* (GLME) with the *lme4* package (Bates, Maechler, Bolker, & Walker, 2018). We tried to fit these GLME models by incrementally adding factors; in total we created five models: (a) the null p_2 with only responder and asker random effects, (b) model 1 which added responder and asker reputation scores as fixed effects, (c) model 2 which added responder and asker roles as fixed effects, (d) model 3 which added responder and asker time on the CSEd SE as fixed effects, and (e) model 4 which added responders’ questions answered and askers’ questions asked as fixed effects.

For each model, we calculated the *standard error of the mean* (SEM) of the conditional modes of the random-effects coefficients from the fitted model; SEM equals the standard deviation divided by the square root of the sample size. We also used the R package *sjstats* (Lüdtke, 2018b) to report the *intraclass-correlation coefficient* (ICC) scores for responders and askers; the ICC represents the percent of variability due to the group. Finally, we compared the *Akaike information criterion* (AIC) and the *Bayesian information criterion* (BIC), keeping in mind that smaller values are better.

We present SEM, ICC, AIC, and BIC scores for all five p_2 selection models in a summary table at the end. Finally, for each selection model, we used the *sjPlot* R package (Lüdtke, 2018a) to create plots showing the conditional modes of the random-effects coefficients and standard errors by contributor, in other words, the difference between the average predicted response for a given set of fixed-effect values (treatment) and the response predicted for a particular individual. This is a method to see how much any individual differs from the population. We examined these plots to see how they changed as we added factors to the selection model.

Null p_2 Selection Model: Responder and Asker Random-Effects

We started by building the null p_2 model for predicting ties (i.e., answering someone else’s question on the CSEd SE). This model only considered responder and asker random effects, which would include in-degree and out-degree.

p_2 Selection Model 1: Reputation

Our first mixed effects p_2 model for predicting ties added the $\log()$ of *reputation scores* as fixed effects to the responder and asker random effects from the null p_2 .

p_2 Selection Model 2: Role

Our second mixed effects p_2 model for predicting ties added responder’s and asker’s role (i.e., moderator, registered user, or unregistered user) as fixed effects to the previous model.

p_2 Selection Model 3: Time in the CSEd SE

Our third mixed effects p_2 model for predicting ties added the $\log()$ of *time on CSEd SE* as fixed effects to the previous model.

p_2 Selection Model 4: Questions Answered or Asked

Our fourth mixed effects p_2 model for predicting ties added the $\log()$ of *questions answered* by responders and *questions asked* by askers as fixed effects to the previous model.

p_2 Summary Table

Finally, we compiled the statistics from the five selection model and added them to a table for comparison (see below). We noted that the standard errors, AIC, and BIC all generally decreased as we iteratively added features to the models.

##		Model 0	Model 1	Model 2	Model 3	Model 4
##	Responder SE	1.3898	0.0350	0.0343	0.0345	0.0000
##	Responder ICC	0.3118	0.0619	0.0646	0.0642	0.0000
##	Asker SE	1.0810	0.3254	0.2665	0.2728	0.0233
##	Asker ICC	0.3583	0.4492	0.4155	0.4126	0.0441
##	AIC	8454.4358	8168.3302	8148.6612	8149.8042	7200.2724
##	BIC	8481.7302	8213.8209	8230.5446	8249.8838	7307.9446

Discussion

Future Research

In future work, we anticipate adding other variables to the selection model, such as users' geographical location or profession. Also, the variables in the current model reflect prior research on teacher interactions where things such as 'reputation', years of service (in CSEd SE it is 'membership duration', formally designated leaders (in CSEd SE it is 'editor/moderator'), previous ties, etc mattered for selection model.

Account for prior relationship

Add aesthetics to network visualization, such as making the node size reflect the $\log()$ of a user's reputation score.

Check clusters with KliqueFinder

References

- Bates, D., Maechler, M., Bolker, B., & Walker, S. (2018). lme4: Linear mixed-effects models using 'Eigen' and S4 (Version 1.1-19) [R package]. Retrieved from <https://cran.r-project.org/package=lme4>
- Csárdi, G. (2018). igraph: Network analysis and visualization (Version 1.2.2) [R package]. Retrieved from <https://CRAN.R-project.org/package=igraph>
- Lüdtke, D. (2018). sjPlot: Data visualization for statistics in social science (Version 2.6.1) [R package]. Retrieved from <https://cran.r-project.org/package=sjPlot>
- Lüdtke, D. (2018). sjstats: Collection of convenient functions for common statistical computations (Version 0.17.2) [R package]. Retrieved from <https://cran.r-project.org/package=sjstats>
- Macià, M., & García, I. (2016). Informal online communities and networks as a source of teacher professional development: A review. *Teaching and Teacher Education*, 55, 291-307.
- Pedersen, T. L. (2018). ggraph: An implementation of grammar of graphics for graphs and networks (Version 1.0.2) [R package]. Retrieved from <https://CRAN.R-project.org/package=ggraph>
- R Core Team. (2018). R: A language and environment for statistical computing (Version 3.5.0) [Computer software]. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from <https://www.R-project.org/>
- Robinson, D. (2015). stackr: An R package for connecting to the Stack Exchange API [R package]. Retrieved from <https://github.com/dgrtwo/stackr>
- Wenger, E. (1998). *Communities of practice: Learning, meaning, and identity*. New York, NY: Cambridge University Press.

Wenger, E. (2011). *Communities of practice: A brief introduction*.

Wickham, H., Chang, W., & RStudio. (2016). ggplot2: Create elegant data visualisations using the grammar of graphics (Version 2.2.1) [R package]. Retrieved from <https://CRAN.R-project.org/package=ggplot2>

Wickham, H., François, R., Henry, L., Müller, K., & RStudio. (2018). dplyr: A grammar of data manipulation (Version 0.7.6) [R package]. Retrieved from <https://CRAN.R-project.org/package=dplyr>